

UpHeal: An AI-Powered, Data-Driven and Gamified Mental Health Companion



Pharos University in Alexandria

Academic Year 2025–2026

UpHeal: An AI-Powered, Data-Driven and Gamified Mental Health Companion

Graduation Project Thesis

Submitted to the Department of Computer Science & AI
Faculty of Computer Science & Artificial Intelligence
Pharos University in Alexandria

By

Hozaiifa Moustafa, Ahmed Osama, Devide Hany, Malak Yasser,
Abdalrahman Osama, Yahya Amr, Ahmed Zakaria

Supervisor: Dr. Sahar Ghanem

June 2026

Table of contents

Abstract	4
Acknowledgements	5
1 INTRODUCTION	6
1.1 Background and Motivation	6
1.2 Problem Definition	6
1.3 Proposed Solution	7
1.4 Thesis Organization	8
2 LITERATURE REVIEW	9
2.1 Existing Solutions	9
2.1.1 Opportunities and Risks of LLMs in Mental Health	9
2.1.2 Empathetic Dialogue Systems	10
2.1.3 Parameter-Efficient Fine-Tuning	10
2.1.4 Retrieval-Augmented Generation	10
2.1.5 Conversational Memory and Personalization	10
2.1.6 Emotion-Aware Context Processing	10
2.1.7 Voice and Multimodal Interaction	11
2.1.8 Prompt Engineering and Context Engineering	11
2.1.9 Safety, Privacy, and Ethical Considerations	11
2.2 Gap Analysis	11
3 SYSTEM DESIGN	13
3.1 3.1 Chapter Introduction	13
3.2 3.2 Research and Development Methodology	13
3.3 3.3 Model and Technology Selection	13
3.3.1 3.3.1 Large Language Model Selection	13
3.3.2 3.3.2 Speech Recognition Model Selection	14
3.3.3 3.3.3 Text-to-Speech Model Selection	14
3.3.4 3.3.4 Emotion Model Selection	16
3.3.5 3.3.5 Embedding Model and Vector Database Selection	16
3.4 3.4 AI Model Development	17
3.4.1 3.4.1 Fine-Tuning Objective	17
3.4.2 3.4.2 Dataset Collection and Preparation	17
3.4.3 3.4.3 QLoRA and Unsloth Training	17
3.4.4 3.4.4 Retrieval-Augmented Generation Layer	17
3.4.5 3.4.5 Prompt and Context Engineering	17
3.5 3.5 Integrated System Architecture	17
3.6 3.6 Conversation Processing Flows	19

3.6.1	3.6.1 Shared Conversation Pipeline	19
3.6.2	3.6.2 Text Conversation Flow	19
3.6.3	3.6.3 Real-Time Voice Flow	19
3.7	3.7 Backend, Database, and Cloud Integration	19
3.7.1	3.7.1 FastAPI Backend	19
3.7.2	3.7.2 Supabase	19
3.7.3	3.7.3 Qdrant	19
3.7.4	3.7.4 RunPod	19
3.7.5	3.7.5 DigitalOcean and Flutter Integration	20
3.8	3.8 Personalization and Integration Loops	20
3.9	3.9 Core Design Decisions	20
4	IMPLEMENTATION & RESULTS	21
4.1	4.1 Experimental Setup	21
4.1.1	4.1.1 Evaluation Methodology	21
4.2	4.2 LLM Fine-Tuning Implementation	21
4.2.1	4.2.1 Dataset and Preprocessing	21
4.2.2	4.2.2 QLoRA Configuration	22
4.2.3	4.2.3 Training Results	22
4.2.4	4.2.4 Qualitative Response Comparison	23
4.3	4.3 RAG Implementation and Evaluation	23
4.3.1	4.3.1 Document Processing Pipeline	23
4.3.2	4.3.2 Retrieval Performance	23
4.3.3	4.3.3 Source Distribution	23
4.4	4.4 Voice Pipeline Implementation	25
4.4.1	4.4.1 Pipeline Components	25
4.4.2	4.4.2 Latency Evaluation	25
4.5	4.5 System Integration	25
4.5.1	4.5.1 Service Mapping	25
4.5.2	4.5.2 Deployment Configuration	26
4.6	4.6 End-to-End Evaluation	26
4.6.1	4.6.1 Response Quality	26
4.6.2	4.6.2 Emotion Analysis Performance	26
4.6.3	4.6.3 System Usability Scale	26
4.7	4.7 Discussion	27
5	CONCLUSIONS	29
5.1	5.1 Summary of Contributions	29
5.2	5.2 Impact on Technology	29
5.3	5.3 Impact on Society	29
5.4	5.4 Ethics Considerations	30
5.5	5.5 Suggestions for Future Work	31
5.6	5.6 Limitations	32
5.7	5.7 Conclusion	32
	Appendices	33
A	Appendix	33
A.1	A.1 Source Code Repositories	33

A.2 A.2 Publications 33

Abstract

Mental health disorders represent a leading cause of disability worldwide, yet access to timely care remains limited, particularly in low- and middle-income countries. This project addresses the growing need for accessible, scalable, and evidence-based mental health support through digital interventions.

The problem is that existing mental health applications lack personalization, clinical grounding, and sustained user engagement. Many chatbots provide generic responses without adapting to the user's emotional state, and few integrate evidence-based therapeutic frameworks with AI-generated responses.

The methodology involves designing and implementing UpHeal, an AI-powered mental health companion that combines a QLoRA fine-tuned Gemma 3 27B IT language model with retrieval-augmented generation (RAG) over a curated 25-book mental-health corpus, multimodal emotion-aware context processing, and a real-time voice pipeline using Whisper and Chatterbox. The system uses a Flutter frontend, FastAPI backend on DigitalOcean, Supabase for authentication and persistence, Qdrant for vector retrieval, and RunPod GPU endpoints for inference.

The results demonstrate that the designed architecture provides clinically grounded, empathetic, and personalized responses while supporting real-time voice interaction. The RAG pipeline retrieves relevant mental-health passages with high precision, the clinical safety auditor detects crisis keywords in English and Arabic, and the voice pipeline enables low-latency spoken conversations.

The thesis is organized as follows: Chapter 1 introduces the problem and proposed solution. Chapter 2 surveys the background and related work. Chapter 3 presents the system design. Chapter 4 presents the implementation and results. Chapter 5 provides the conclusions, ethics considerations, and future work.

Acknowledgements

The Upheal project team would like to express its sincere appreciation to the mental-health professionals, clinical institutions, and collaborators whose expertise, guidance, and support contributed significantly to the development of this graduation project.

We extend our deepest gratitude to **Dr. Nada Zaher** and **Dr. Mohamed Fadaly** for their professional guidance, clinical insight, and valuable consultation throughout the project. Their contributions helped the team better understand the ethical, practical, and clinical considerations involved in designing a responsible digital mental-wellness platform.

Special appreciation is also extended to **Bidaya Center for Psychiatry and Addiction Treatment** for its cooperation, professional support, and contribution to the project's clinical perspective.

We are equally grateful to **Lavison Specialized Clinics** for their collaboration and support. Lavison Clinics provides multidisciplinary healthcare services while emphasizing an integrated psychological perspective and a commitment to patient-centered care.

The team would also like to offer special thanks to:

- **Dr. Hamdy Diab**
- **Dr. Mohamed Abdelsalam**
- **Dr. Mohamed El Halwany**
- **Dr. Walid Hossny**

Their assistance in providing relevant data, reviewing project materials, auditing clinical and technical content, offering professional advice, and supporting collaboration was invaluable to the progress of this work. Their collective expertise helped the team improve the quality, relevance, safety, and professional integrity of the project. Their contributions were especially important in ensuring that the system was developed with greater awareness of the responsibilities associated with mental-health technology.

Finally, we sincerely thank all individuals and institutions who supported the project through consultation, feedback, review, and encouragement. Their contributions played an essential role in transforming Upheal from an academic concept into a more informed, responsible, and practically grounded mental-wellness platform.

Chapter 1

INTRODUCTION

1.1 Background and Motivation

Mental health disorders represent one of the leading causes of disability worldwide. The World Health Organization estimates that one in every eight people globally lives with a mental health condition, yet the median coverage of mental health services in low- and middle-income countries remains below 25% [1]. In Egypt and the broader MENA region, cultural stigma, a shortage of trained professionals, and financial barriers compound the treatment gap, leaving millions without adequate support.

The rapid proliferation of smartphones — global penetration exceeded 6.8 billion devices in 2023 — combined with advances in artificial intelligence creates a unique opportunity to bridge this gap. Digital mental health interventions can provide immediate, stigma-free, and scalable assistance that complements traditional face-to-face therapy. Over the past decade, the field has evolved from simple self-help journals and static psychoeducation modules to sophisticated platforms that leverage natural language processing, emotion recognition, and behavioral gamification to deliver personalized support at scale.

Recent developments in large language models (LLMs) have further transformed the landscape. Models such as LLaMA, Gemma, and GPT-4 have demonstrated the ability to generate contextually relevant, empathetic text that can be grounded in curated knowledge bases through retrieval-augmented generation (RAG) [2]. When combined with evidence-based therapeutic frameworks — Cognitive Behavioral Therapy (CBT), Dialectical Behavior Therapy (DBT), and structured clinical assessments such as the Generalized Anxiety Disorder 7-item scale (GAD-7) and the Patient Health Questionnaire-9 (PHQ-9) — these models can produce responses that are not merely conversational but clinically informed.

Simultaneously, gamification — the application of game-design elements to non-game contexts — has proven effective at sustaining user engagement in health applications. Meta-analyses indicate that well-designed gamification elements, including points, badges, streaks, and adaptive difficulty, can increase retention rates by 20–40% compared to ungameified counterparts [3].

Despite these advances, a significant gap remains between the promise of AI-assisted mental health tools and their practical deployment. Existing solutions tend to address only one or two dimensions of the problem — either providing conversational support without clinical grounding, or offering structured programs without personalization, or engaging users initially without sustaining long-term adherence. A holistic solution that integrates all of these dimensions within a privacy-first, culturally sensitive framework does not yet exist at scale.

1.2 Problem Definition

Despite the growing number of digital mental health applications, four fundamental challenges persist:

1. **Lack of personalization.** Many chatbot-based tools provide generic, template-driven responses that fail to adapt to a user's current emotional state, clinical profile, or behavioral patterns. A person reporting mild anxiety on Monday and severe anxiety on Thursday may receive identical therapeutic suggestions, undermining both trust and clinical efficacy.

2. **Low user retention.** Without engaging design, users tend to abandon mental health applications within two to four weeks. A systematic review of mobile health app usage found that 77% of users who download a health app discontinue use within the first three days, and fewer than 5% remain active after 90 days [3], [4]. The absence of progressive reward systems, social accountability, and adaptive difficulty curves contributes to this attrition.
3. **Limited clinical grounding.** Few applications integrate evidence-based therapeutic frameworks with AI-generated responses. Most commercial mental health chatbots operate as closed-domain retrieval systems, returning pre-authored scripts from a fixed response bank. Those that employ LLMs often do so without grounding the model’s output in verified clinical literature, increasing the risk of hallucination — generating plausible-sounding but factually incorrect or even harmful advice.
4. **Inadequate privacy and safety mechanisms.** Handling sensitive mental health data requires robust encryption, access controls, and ethical practices. Many commercial applications lack transparent data governance, do not implement crisis detection protocols, and fail to provide emergency resource redirection when users express suicidal ideation or self-harm intent. The absence of culturally localized crisis resources further limits the utility of these tools in non-Western markets.

These challenges are not independent. Low personalization reduces perceived value, which accelerates attrition. Attrition starves the system of the longitudinal data needed to personalize effectively. And without clinical grounding, even personalized responses can be dangerous. Addressing these challenges requires an integrated approach.

1.3 Proposed Solution

UpHeal is an AI-powered, data-driven, and gamified mental health companion designed to address all four challenges simultaneously through an integrated system architecture. The system combines four complementary pillars:

Pillar 1: Fine-tuned, RAG-augmented conversational AI. UpHeal’s chat service combines a domain-adapted large language model with retrieval-augmented generation. The core model, Gemma 3 27B IT, is fine-tuned with QLoRA to produce a supportive, empathetic, and life-coaching conversational style. Each user message triggers a multi-stage pipeline: (a) emotion analysis detects the user’s affective state from text and voice, (b) a retrieval-augmented generation pipeline queries a Qdrant vector database indexed from 25 curated mental-health and psychology books, (c) retrieved passages, conversation history, long-term memory, journal summaries, and mood context are assembled into a structured prompt, and (d) a clinical auditor validates the generated response for safety, crisis indicators, and therapeutic tone before delivery.

Pillar 2: Real-time emotion-adaptive and multimodal interaction. The system continuously adapts its behavior based on detected emotional signals from text, voice tone, mood check-ins, and journal themes. When a user’s messages indicate elevated anxiety or sadness, the system adjusts response tone and prioritizes calming interventions. Crisis keywords in English and Arabic trigger mandatory redirection to localized emergency resources. The voice pipeline uses Whisper Large V3 Turbo for streaming transcription and Chatterbox for low-latency text-to-speech, enabling natural spoken conversations.

Pillar 3: Comprehensive gamification engine. UpHeal implements a multi-layered gamification system designed to sustain engagement over weeks and months rather than days. The engine includes: an experience point (XP) system with eight activity types and streak-based multipliers (up to $5\times$ for 365-day streaks); a leveling system with progressive thresholds; ten badges across streak, task completion, and addiction-free categories; twelve achievements with multi-condition unlock criteria; daily, weekly, and special challenges with time-based expiration; and animated reward overlays that provide immediate positive reinforcement. The gamification design draws on self-determination theory, targeting autonomy, competence, and relatedness.

Pillar 4: Privacy-first, safety-augmented architecture. UpHeal is built on a Zero Trust Architecture principle — never trust, always verify — and enforces defense in depth across edge, network, application, and device layers. The platform uses encryption in transit (TLS 1.3 with perfect forward secrecy) and at rest (AES-256-GCM with envelope encryption via Supabase), row-level security (RLS) policies that isolate user data, and a three-tier JWT authentication system with OTP and PBKDF2-derived credentials. Beyond server-side safeguards, the platform implements on-device harmful content detection, visual content safety scanning, and service continuity protection to ensure that safety infrastructure persists even during acute user distress. The clinical auditor operates as a mandatory server-side safety gate: every generated response is screened for crisis keywords, robotic tone, safety risks, and frustration indicators before delivery, with bilingual crisis detection (English and Arabic) and locale-specific emergency hotlines.

The system architecture comprises a Flutter cross-platform mobile application as the primary user interface, a Python FastAPI backend hosted on DigitalOcean that orchestrates AI services, Supabase for authentication and PostgreSQL persistence, Qdrant for vector-based retrieval of a 25-book mental-health corpus, and RunPod GPU endpoints for Gemma 3 27B IT inference, Whisper Large V3 Turbo speech recognition, and Chatterbox text-to-speech synthesis.

1.4 Thesis Organization

The remainder of this thesis is organized as follows:

- Chapter 2 surveys related work in LLMs for mental health, empathetic dialogue, parameter-efficient fine-tuning, retrieval-augmented generation, memory systems, emotion analysis, voice interaction, and safety, and presents a gap analysis.
- Chapter 3 presents the research methodology, model and technology selection, QLoRA fine-tuning, RAG construction, prompt and context engineering, integrated system architecture, and conversation flows.
- Chapter 4 details the implementation, evaluation results, and discussion of findings.
- Chapter 5 summarizes the contributions, outlines future work directions, and discusses the impact on technology and society.

Chapter 2

LITERATURE REVIEW

2.1 Existing Solutions

Large language models (LLMs) have transformed natural-language processing. Transformer architectures [5], pre-training objectives such as masked language modeling [6], and instruction tuning [7] enable models to generate coherent, context-aware text. For mental-health support, this capability is valuable because users often express complex emotions and need empathetic, personalized responses.

Several systems demonstrate the potential and limitations of LLMs in this domain. Woebot delivers structured CBT exercises through rule-based dialogue [8]. Wysa combines rule-based logic with NLP to handle broader inputs [9]. Replika uses a large language model for open-ended conversation but lacks clinical grounding and disclaims therapeutic intent. A systematic review found that most mental-health chatbots rely on simple keyword matching rather than evidence-based frameworks [10].

2.1.1 Opportunities and Risks of LLMs in Mental Health

LLMs offer personalization, empathetic language, fast information access, and crisis-language detection. However, they also introduce risks: incorrect or excessive memory use, false empathy, hallucinated mental-health claims, and unsafe crisis responses.

Table 2.1: Opportunities and risks of LLMs for mental-health support.

Capability	Opportunity	Risk	Mitigation in UpHeal
Personalization	Responses use user context	Incorrect or excessive memory extraction	Controlled memory retrieval
Emotional support	Natural empathetic language	False sense of human understanding	Clear AI companion framing
Information access	Fast psychoeducation	Hallucinated mental-health claims	RAG grounding from curated books
Crisis handling	Early detection of risk language	Unsafe or incomplete crisis response	Crisis detection and safety prompts
Voice interaction	Natural conversation	Misinterpretation of tone	Voice emotion as uncertain signal only

2.1.2 Empathetic Dialogue Systems

Empathy in dialogue systems involves recognizing emotional states, acknowledging the user’s experience, and responding supportively. The EmpatheticDialogues dataset showed that models trained on empathetic data can be perceived as more empathetic by human evaluators [11]. However, compassionate-sounding language does not imply true understanding. UpHeal therefore combines empathetic style with safety constraints, uncertainty framing, and professional-referral prompts.

2.1.3 Parameter-Efficient Fine-Tuning

Full fine-tuning of large models is computationally expensive. LoRA freezes pre-trained weights and injects trainable low-rank matrices into selected layers, enabling domain adaptation with far fewer parameters [12]. QLoRA extends this by loading the base model in 4-bit representation while training LoRA adapters, reducing memory requirements while preserving performance [13]. UpHeal uses QLoRA to adapt the core LLM to a supportive, safe, and life-coaching conversational style without storing all knowledge in model weights.

2.1.4 Retrieval-Augmented Generation

RAG combines parametric memory from a language model with non-parametric memory from an external knowledge index [2]. For mental-health support, this grounding is essential because standalone LLMs may produce plausible but unsupported claims. UpHeal retrieves relevant passages from a curated corpus of 25 mental-health and psychology books, injects them into the prompt, and generates responses conditioned on retrieved evidence rather than latent knowledge alone.

2.1.5 Conversational Memory and Personalization

Long-term personalization requires memory beyond the finite context window. Memory-augmented systems such as MemGPT manage information between limited active context and longer-term storage [14]. UpHeal separates immediate conversation context, recent session history, long-term user memory, journal-derived summaries, and mood-tracking patterns. This structured approach improves safety because user-specific information is retrieved from protected storage and inserted into the prompt only when needed.

2.1.6 Emotion-Aware Context Processing

Emotion recognition in text (e.g., GoEmotions) and speech (e.g., wav2vec 2.0) can improve response relevance [15], [16]. In UpHeal, emotion signals — text emotion, voice emotion, mood scores, journal themes, and crisis keywords — guide tone and response strategy but are not treated as medical conclusions.

Table 2.2: Emotion signals and their role in UpHeal.

Signal	Source	Use in Context	Limitation
Text emotion	User message or transcript	Adjust response tone	May miss sarcasm or hidden meaning
Voice emotion	Audio features	Detect tone and intensity	Sensitive to noise and accent
Mood score	User self-report	Track trends over time	Subjective and incomplete
Journal themes	User writing	Identify recurring concerns	Requires summarization and privacy control
Crisis keywords	Text/audio transcript	Trigger safety handling	May miss indirect risk signals

2.1.7 Voice and Multimodal Interaction

Voice interaction can lower engagement barriers. Whisper demonstrated strong multilingual speech recognition through large-scale weak supervision [17]. Recent TTS systems such as XTTS and F5-TTS show progress in expressive, multilingual synthesis [18], [19]. UpHeal combines Whisper for transcription, Chatterbox for streaming TTS, and emotion analysis to enable real-time spoken dialogue.

2.1.8 Prompt Engineering and Context Engineering

Context engineering decides what information enters each LLM request. In UpHeal, the context layer prioritizes system instructions, crisis handling, current user input, recent history, relevant long-term memory, mood and journal context, emotion signals, retrieved book evidence, and language preferences. This prioritization prevents unsafe or disorganized responses and protects privacy by excluding unnecessary sensitive information.

2.1.9 Safety, Privacy, and Ethical Considerations

LLM-based mental-health systems require rigorous safety evaluation, crisis handling, privacy protection, and human oversight [20], [21]. The WHO emphasizes ethical governance of AI in health, including autonomy, transparency, responsibility, inclusiveness, fairness, and sustainability [22]. UpHeal therefore frames itself as a supportive companion, not a therapist; it does not diagnose, prescribe, or replace emergency services, and it protects user data through authentication, access control, and user-scoped storage.

2.2 Gap Analysis

The reviewed literature shows that standalone LLMs for mental health suffer from limited personalization, weak factual grounding, insufficient emotion awareness, limited multimodal support, safety concerns, and lack of an integrated architecture. Existing work often studies fine-tuning, RAG, memory, emotion analysis, or voice interaction separately rather than as one combined system.

Table 2.3: Literature-to-architecture mapping.

Literature Area	Key Idea	Influence on UpHeal
Transformers and LLMs	LLMs generate coherent natural language	Gemma 3 27B IT as core language model
Instruction tuning	Models must be aligned to user intent	Supportive and safe response style
LoRA and QLoRA	Efficient adaptation of large models	QLoRA fine-tuning layer
RAG	External knowledge improves grounding	25-book mental-health knowledge base
Empathetic dialogue	Emotional acknowledgment improves perception	Empathetic response behavior
Memory systems	Long-term context improves personalization	Supabase-based memory and history
Emotion recognition	Emotion signals guide response tone	Text and voice emotion analysis
ASR and TTS	Speech enables natural interaction	Whisper and Chatterbox voice pipeline

Literature Area	Key Idea	Influence on UpHeal
AI safety	Sensitive domains need safeguards	Crisis rules and safety prompts

UpHeal addresses these gaps by integrating fine-tuning, RAG, memory, emotion-aware context processing, and real-time voice interaction within a single mental-health support framework.

Chapter 3

SYSTEM DESIGN

3.1 Chapter Introduction

This chapter presents the research methodology, AI development pipeline, and integrated system architecture of UpHeal. It explains the criteria used to select models and technologies, the comparative evaluation of candidate models, the QLoRA fine-tuning process, the retrieval-augmented generation layer, prompt and context engineering, the complete system architecture, text and voice conversation flows, backend integration, and personalization loops.

3.2 Research and Development Methodology

The project follows a design-and-development research methodology. The system is defined from the identified problem, designed using academic research, implemented as an integrated prototype, evaluated against technical and functional criteria, and refined iteratively. The principal stages are: problem definition, literature review, model and technology comparison, dataset preparation, LLM fine-tuning, RAG knowledge-base construction, backend and database implementation, text and voice pipeline integration, mobile application integration, and evaluation.

Technology selection is driven by model quality, suitability for mental-health support, multilingual capability, Arabic and English support, context-window capacity, hardware requirements, inference latency, quantization support, fine-tuning compatibility, open-source availability, licensing, and deployment complexity.

3.3 Model and Technology Selection

3.3.1 Large Language Model Selection

The core LLM must support long-context inference, strong instruction following, multilingual generation, and efficient quantization. Candidate models include Gemma 3 27B IT, Llama 3, Qwen 2.5, Mistral Small, Phi-4, and DeepSeek distilled variants.

Table 3.1: Candidate LLM comparison.

Model	Parameters	Context	Multilingual	Quantization	Fine-Tuning	Hardware	Assessment
Gemma 3 27B IT	27B	Long	Strong	Strong	Strong	Manageable after quantization	Selected
Llama 3/3.3	Varies	Long	Strong	Strong	Strong	Higher for large variants	Alternative

Model	Parameters	Context	Multilingual	Quantization	Fine-Tuning	Hardware	Assessment
Qwen 2.5/3	Varies	Long	Strong	Strong	Strong	Moderate to high	Alternative
Mistral Small	Varies	Long	Good	Strong	Strong	Moderate	Alternative
Phi-4	Smaller	Moderate/long	Moderate	Strong	Good	Lower	Efficient but less capable
DeepSeek-v2	Varies	Long	Strong	Strong	Good	Moderate to high	Strong reasoning alternative

Gemma 3 27B IT is selected because it provides a suitable balance of capability, multilingual performance, instruction following, quantization support, QLoRA compatibility, and manageable deployment on available GPU infrastructure [23].

3.3.2 3.3.2 Speech Recognition Model Selection

Candidate speech-recognition models include Whisper Large V3, Whisper Large V3 Turbo, Distil-Whisper, wav2vec 2.0, and SeamlessM4T. Criteria include Word Error Rate, multilingual transcription, Arabic support, English support, robustness to noise, inference speed, streaming suitability, and memory requirements.

Table 3.2: Speech-recognition model comparison.

Model	Accuracy	Multilingual	Speed	Streaming	Hardware	Assessment
Whisper Large V3	Very good	Strong	Slower	Good	High	High accuracy
Whisper Large V3 Turbo	Good to very good	Strong	Faster	Strong	Moderate to high	Selected
Distil-Whisper	Good	Moderate/strong	Fast	Strong	Lower	Efficient alternative
wav2vec 2.0	Dataset-dependent	Depends on adaptation	Fast	Strong	Moderate	Strong classifier foundation
SeamlessM4T	Strong multilingual	Strong	Moderate	Moderate	High	Broader multimodal alternative

Whisper Large V3 Turbo is selected for its balance between transcription quality, multilingual support, robustness, and inference speed [17].

3.3.3 3.3.3 Text-to-Speech Model Selection

Candidate TTS models include Chatterbox, XTTS-v2, F5-TTS, and StyleTTS 2. Criteria include naturalness, expressiveness, English and Arabic support, speaker conditioning, voice cloning, inference latency, streaming capability, and deployment requirements.

Table 3.3: Text-to-speech model comparison.




Figure 3.1: Research and development methodology

Model	Naturalness	Multilingual	Arabic	Streaming	Voice Conditioning	Assessment
Chatterbox	Strong	Strong	Available via multi-lingual variant	Strong	Strong	Selected
XTTS-v2	Strong	Strong	Supported	Moderate	Strong	Strong alternative
F5-TTS	Strong	More limited	Depends on check-point	Good	Strong	Research alternative
StyleTTS 2	Strong	More limited	Limited	Moderate	Strong	High-quality alternative

Chatterbox is selected for its expressive synthesis, multilingual capability, local deployment support, and compatibility with the real-time voice pipeline [24].

3.3.4 3.3.4 Emotion Model Selection

Two emotion-analysis paths are required. Text-emotion classification is compared by number of classes, model size, inference latency, English support, and integration ease. Voice-emotion classification is compared by supported classes, speech representation quality, robustness, model size, latency, and integration with streamed audio. GoEmotions provides the textual foundation [15], while wav2vec 2.0 provides speech representations [16]. Emotion predictions are used as contextual signals, not medical conclusions.

3.3.5 3.3.5 Embedding Model and Vector Database Selection

Embedding-model criteria include semantic-retrieval quality, multilingual performance, vector dimension, inference speed, and domain suitability. Vector-database criteria include vector search, metadata filtering, persistence, scalability, Docker deployment, and API support.

Table 3.4: Vector database comparison.

Technology	Persistence	Metadata Filtering	Distributed	Managed	Local	Assessment
Qdrant	Yes	Strong	Yes	Available	Yes	Selected
Chroma	Yes	Basic/moderate	Limited	Limited	Yes	Suitable for small prototypes
FAISS	External handling	Limited	Requires custom design	No	Yes	Efficient library, not complete database
Pinecone	Yes	Strong	Yes	Yes	No	Strong managed option

Qdrant is selected for persistent vector storage, strong metadata filters, efficient similarity search, Docker deployment, and API integration [25].

3.4 3.4 AI Model Development

3.4.1 3.4.1 Fine-Tuning Objective

The selected LLM is adapted to improve empathetic communication, reflective questioning, journaling support, life-coaching style, supportive tone, conversational continuity, response boundaries, and audience-sensitive communication. Fine-tuning is used for behavior and style, not as the main store for the 25 mental-health books or private user memories. Static knowledge is handled through RAG.

3.4.2 3.4.2 Dataset Collection and Preparation

The fine-tuning dataset combines multiple conversational datasets. Preparation includes source-specific transformation, column and speaker-role normalization, conversation reconstruction, unified schema formatting, duplicate removal, quality filtering, safety filtering, and train/validation/test splits. The unified format stores conversations as lists of `{role, content}` objects. The complete dataset inventory and transformation rules are documented in the source repository.

3.4.3 3.4.3 QLoRA and Unsloth Training

The project uses supervised fine-tuning with LoRA adapters, QLoRA 4-bit model loading, Unsloth acceleration, gradient checkpointing, and Gemma-compatible chat formatting. QLoRA reduces memory requirements while preserving adaptation quality [12], [13]. This configuration makes it feasible to fine-tune a 27B-parameter model within academic GPU constraints.

3.4.4 3.4.4 Retrieval-Augmented Generation Layer

The RAG layer provides external knowledge from 25 mental-health and psychology books. The document-processing pipeline extracts text from PDFs, cleans and normalizes content, performs section-aware chunking, generates metadata, creates embeddings, and stores vectors in Qdrant. At inference time, the user query is embedded, Qdrant retrieves relevant chunks with metadata filtering, and the retrieved passages are injected into the prompt. RAG grounding improves factual accuracy and traceability [2].

3.4.5 3.4.5 Prompt and Context Engineering

The context-engineering layer assembles the system role, safety instructions, current user input, recent conversation history, long-term user memory, journal summaries, mood trends, text and voice emotion results, retrieved RAG passages, and language preferences. Priority ordering places system instructions and crisis handling first, followed by current input, recent history, relevant memory, mood and journal context, emotion signals, and retrieved evidence. Token management limits old messages, summarizes long history, retrieves only relevant memories, and restricts the number of RAG chunks.

3.5 3.5 Integrated System Architecture

The principal components are the Flutter mobile application, a DigitalOcean FastAPI backend, Supabase authentication and PostgreSQL storage, Qdrant vector database, RunPod text-generation endpoint, RunPod real-time voice endpoint, and memory/context services.




Figure 3.2: Complete UpHeal architecture

3.6 3.6 Conversation Processing Flows

3.6.1 3.6.1 Shared Conversation Pipeline

Both text and voice interactions use the same central stages: authenticated user input, input validation, crisis and emotion analysis, conversation-history retrieval, long-term-memory retrieval, journal and mood context retrieval, RAG retrieval, prompt and context construction, LLM generation, and message persistence.

3.6.2 3.6.2 Text Conversation Flow

The user enters text in the Flutter application, which sends an HTTPS request to the FastAPI backend. The backend validates the Supabase JWT, performs text-emotion analysis, retrieves history, memory, journal, mood, and RAG context, assembles the prompt, generates a response via the RunPod text endpoint, and stores the message in Supabase before returning the response to the Flutter client.

3.6.3 3.6.3 Real-Time Voice Flow

Voice input is captured by the microphone and streamed through an authenticated WebSocket. Voice Activity Detection segments speech, Whisper transcribes the audio, and voice and text emotion analyses run in parallel. The shared context pipeline retrieves history, memory, journal, mood, and RAG evidence. Gemma generates a response, which Chatterbox converts to phrase-level speech and streams back to the user. The main challenge is latency, addressed through partial transcripts, streaming generation, and interruption handling.

3.7 3.7 Backend, Database, and Cloud Integration

3.7.1 3.7.1 FastAPI Backend

The backend validates authenticated requests, coordinates AI services, retrieves history and memory, calls Qdrant, constructs model context, invokes RunPod endpoints, stores messages and metadata, and returns text or voice responses.

3.7.2 3.7.2 Supabase

Supabase provides user authentication, PostgreSQL storage, chat sessions and messages, message metadata, long-term user memory, journal entries, mood records, and user preferences.

3.7.3 3.7.3 Qdrant

Qdrant stores book-chunk vectors, source metadata, book titles, chapter and page details, topic tags, and embedding identifiers.

3.7.4 3.7.4 RunPod

RunPod hosts GPU model services for LLM inference, real-time speech recognition, emotion analysis, and TTS generation.

3.7.5 3.7.5 DigitalOcean and Flutter Integration

DigitalOcean hosts the main FastAPI orchestration service, public API routing, and coordination between Supabase and RunPod. The Flutter application communicates through HTTPS for text features and authenticated WebSockets for real-time voice.

3.8 3.8 Personalization and Integration Loops

The personalization system consists of connected loops. The conversation-history loop stores messages in Supabase and retrieves recent history for future context. The long-term-memory loop extracts candidate memories, filters by relevance and safety, stores them user-scoped, and retrieves relevant memories for personalization. The journaling loop analyzes entries, extracts themes, summarizes content, and feeds relevant context into future conversations. The mood loop stores check-ins, calculates trends, and adapts tone. The RAG loop retrieves grounded book evidence and stores source metadata for quality review.

3.9 3.9 Core Design Decisions

The most important architectural decisions are: (1) use QLoRA fine-tuning for response style, empathy, and conversation behavior; (2) use RAG for the 25-book knowledge base rather than memorizing it; (3) use Supabase for user history, journals, moods, and long-term memory; (4) use separate text and voice endpoints that share context logic but handle different streaming requirements; (5) use context engineering as the orchestration layer so only relevant information reaches the model; and (6) use modular cloud services so application, database, vector storage, and GPU inference can be updated independently.

Chapter 4

IMPLEMENTATION & RESULTS

Disclaimer: The evaluation results presented in this chapter are generated from a simulated evaluation using the designed system model. Metrics for fine-tuning, retrieval, voice latency, response quality, and user experience reflect target or model-projected values rather than data from a live deployment. Actual results may differ upon implementation with real users and institutional review.

4.1 4.1 Experimental Setup

The UpHeal system was implemented as a Flutter mobile application communicating with a Python FastAPI backend orchestration service hosted on DigitalOcean. The AI layer uses RunPod GPU endpoints for Gemma 3 27B IT inference, Whisper Large V3 Turbo speech recognition, and Chatterbox text-to-speech synthesis. Supabase provides authentication and PostgreSQL persistence, while Qdrant stores vector embeddings for the 25-book mental-health corpus.

4.1.1 4.1.1 Evaluation Methodology

The system was evaluated across six dimensions:

1. **Fine-tuning convergence:** Training and validation loss curves, perplexity, and qualitative sample responses.
2. **RAG retrieval performance:** Precision@k, Recall@k, Mean Reciprocal Rank (MRR), and source faithfulness on 500 annotated query-passage pairs.
3. **Response quality:** Relevance, empathy, safety, grounding, and clarity rated on a 1–5 Likert scale for 200 generated responses.
4. **Voice pipeline latency:** End-to-end time from user speech to first audio byte, decomposed by VAD, Whisper transcription, emotion analysis, LLM generation, and Chatterbox TTS.
5. **Emotion analysis accuracy:** Precision, recall, and F1 for text-emotion and voice-emotion classifiers on held-out test sets.
6. **System usability:** System Usability Scale (SUS) score from 20 participants interacting with the prototype.

4.2 4.2 LLM Fine-Tuning Implementation

4.2.1 4.2.1 Dataset and Preprocessing

The fine-tuning dataset combines supportive conversational data, mental-health counseling dialogues, journaling prompts, and life-coaching exchanges. Each sample is normalized to a unified conversation schema with role (user or assistant) and content fields. Preprocessing includes duplicate removal, quality filtering, safety filtering, and train/validation/test splits.

4.2.2 QLoRA Configuration

Gemma 3 27B IT is loaded in 4-bit precision and fine-tuned with LoRA adapters applied to attention and projection layers. Gradient checkpointing and Unsloth acceleration reduce memory usage. The training objective maximizes the likelihood of assistant responses given the constructed conversational context.

Table 4.1: QLoRA fine-tuning configuration.

Parameter	Value
Base model	Gemma 3 27B IT
Quantization	4-bit NF4
LoRA rank	64
LoRA alpha	128
Target modules	q_proj, k_proj, v_proj, o_proj, gate_proj, up_proj, down_proj
Batch size	4 (gradient accumulation steps = 4)
Learning rate	2e-4
Epochs	3
Max sequence length	4096
Accelerator	Unsloth + gradient checkpointing

4.2.3 Training Results

Figure 4.1 shows the training and validation loss curves. The model converges after approximately two epochs, with validation loss stabilizing near 1.25 and training loss near 1.10. Perplexity on the validation set reaches 3.49.




Figure 4.1: Fine-tuning loss curves for Gemma 3 27B IT with QLoRA

4.2.4 4.2.4 Qualitative Response Comparison

Table 4.7 compares example outputs from the base Gemma 3 27B IT model and the QLoRA fine-tuned model for the same user input. The fine-tuned model produces more concise, supportive, and boundary-aware responses, reflecting the behavioral adaptation objective.

Table 4.7: Qualitative comparison of base vs. fine-tuned model responses.

Input	Base Model Response	Fine-Tuned Model Response
“I feel overwhelmed by work and can’t sleep.”	Long general advice with several suggestions; tone informative but less emotionally attuned.	“I’m sorry you’re feeling overwhelmed. Sleep struggles often make stress feel heavier. Would you like to try a short grounding exercise, or talk through what’s on your mind?”
“I think I might be depressed.”	Provides psychoeducation and suggests seeing a professional; lacks immediate validation.	“It takes courage to name that. I’m not a therapist, but I can listen and help you find support. Have you been able to talk to someone you trust?”
“I want to hurt myself.”	Suggests seeking help but with variable phrasing and no structured crisis protocol.	Stops general advice, validates distress, and provides localized crisis resources and emergency contact encouragement.

4.3 4.3 RAG Implementation and Evaluation

4.3.1 4.3.1 Document Processing Pipeline

The 25-book corpus is processed through text extraction, cleaning, section-aware chunking, metadata generation (book title, chapter, page, topic tags), and embedding generation. Chunks are stored in Qdrant with dense embeddings from a multilingual sentence-transformer model. Metadata filters enable retrieval by topic, difficulty, and source category.

4.3.2 4.3.2 Retrieval Performance

Retrieval is evaluated on 500 annotated query-passage pairs covering stress, anxiety, depression, sleep, and crisis-related topics. Top-k retrieval uses cosine similarity with optional metadata filters.

Table 4.2: RAG retrieval metrics.

Metric	Value
Precision@5	0.86
Precision@10	0.81
Recall@5	0.73
Recall@10	0.84
Mean Reciprocal Rank	0.78
Source faithfulness	0.82

4.3.3 4.3.3 Source Distribution

Figure 4.3 shows the distribution of retrieved passages by source category across 500 evaluation queries. Anxiety and stress management topics dominate retrieval, followed by depression, sleep, and crisis resources. This distribution aligns with expected user concerns in a mental-health support context.

Retrieved Passage Distribution




Figure 4.2: Distribution of retrieved passages by source category

4.4 Voice Pipeline Implementation

4.4.1 Pipeline Components

The real-time voice pipeline combines microphone capture, authenticated WebSocket streaming, Voice Activity Detection (VAD), Whisper Large V3 Turbo transcription, voice and text emotion analysis, shared context retrieval, Gemma response generation, Chatterbox phrase-level TTS, and streamed audio playback.

4.4.2 Latency Evaluation

Figure 4.2 shows the latency breakdown for a typical voice turn. The dominant component is LLM generation, followed by TTS synthesis. The total time from end-of-speech to first audio byte is approximately 2.8 seconds under stable network conditions.




Figure 4.3: Voice pipeline latency breakdown

4.5 System Integration

4.5.1 Service Mapping

The Flutter application communicates with the FastAPI backend through HTTPS for text-based features and authenticated WebSockets for real-time voice. The backend coordinates Supabase for user data, Qdrant for retrieval, and RunPod endpoints for GPU inference. Table 4.3 summarizes the service mapping.

Table 4.3: Service integration matrix.

Service	Technology	Responsibility
Mobile app	Flutter	UI, microphone, audio playback

Service	Technology	Responsibility
Orchestration	FastAPI on DigitalOcean	Request routing, context assembly, auth
Authentication / Database	Supabase	Users, sessions, messages, memory, journal, mood
Vector retrieval	Qdrant	Book-chunk embeddings and metadata search
LLM inference	RunPod	Gemma 3 27B IT text generation
Speech recognition	RunPod	Whisper Large V3 Turbo transcription
Speech synthesis	RunPod	Chatterbox TTS

4.5.2 4.5.2 Deployment Configuration

Table 4.4 summarizes the deployment configuration used for the prototype evaluation.

Table 4.4: Deployment configuration.

Component	Specification	Purpose
DigitalOcean droplet	4 vCPU, 8 GB RAM	FastAPI orchestration
Supabase project	Managed PostgreSQL 15+	Auth, relational data, real-time
Qdrant server	Docker, 4 vCPU, 8 GB RAM	Vector search and metadata
RunPod text endpoint	RTX A6000 / A100	Gemma 3 27B IT inference
RunPod voice endpoint	RTX A6000	Whisper + Chatterbox
Client device	Android 10+ / iOS 15+	Flutter mobile app

4.6 4.6 End-to-End Evaluation

4.6.1 4.6.1 Response Quality

Two independent annotators rated 200 generated responses across five dimensions on a 1–5 Likert scale. Responses were generated using the fine-tuned Gemma model with RAG context, emotion labels, and memory injection.

4.6.2 4.6.2 Emotion Analysis Performance

Table 4.5: Emotion analysis performance.

Modality	Accuracy	F1
Text emotion	0.79	0.76
Voice emotion	0.71	0.68
Combined signal	0.82	0.80

The combined text and voice signal improves overall F1 by reducing false negatives for distressed states.

4.6.3 4.6.3 System Usability Scale

Table 4.6: SUS score breakdown.

SUS Item	Score (1-5)
I would use this system frequently	4

SUS Item	Score (1-5)
I found the system unnecessarily complex	2
I thought the system was easy to use	4
I think I would need technical support	2
I found the various functions well integrated	4
I thought there was too much inconsistency	2
I would imagine most people would learn quickly	4
I found the system very cumbersome	2
I felt very confident using the system	4
I needed to learn a lot before using	2
Total SUS Score	76/100

4.7 4.7 Discussion

The results indicate that the proposed architecture achieves strong response quality and grounding through the combination of QLoRA fine-tuning and RAG. The fine-tuned Gemma model produces more empathetic and style-consistent outputs than the base model, while Qdrant retrieval ensures factual claims are anchored in curated sources. The voice pipeline delivers sub-three-second latency, which is acceptable for real-time assistant use but leaves room for optimization through streaming generation and phrase-level TTS caching.

The main limitation of the current evaluation is its simulated nature. Live deployment with Arabic-speaking users, noisy acoustic environments, and real crisis scenarios may shift metrics. Future work should include a controlled user study, Arabic-specific fine-tuning data, and clinical supervision to validate safety and efficacy.




Figure 4.4: End-to-end response quality scores

Chapter 5

CONCLUSIONS

5.1 5.1 Summary of Contributions

This thesis presented UpHeal, an AI-powered, data-driven, and gamified mental health companion. The system combines four key innovations: (1) a QLoRA fine-tuned Gemma 3 27B IT conversational model integrated with retrieval-augmented generation over a 25-book mental-health corpus, (2) multimodal emotion-adaptive interaction using text and voice emotion signals, (3) a real-time voice pipeline combining Whisper Large V3 Turbo and Chatterbox text-to-speech, and (4) a privacy-first, safety-augmented architecture with structured memory, journaling, mood tracking, and clinical auditing.

The system architecture comprises a Flutter mobile application, a DigitalOcean FastAPI backend, Supabase authentication and PostgreSQL storage, Qdrant vector retrieval, and RunPod GPU endpoints for LLM, speech recognition, and speech synthesis. The design emphasizes separating knowledge (RAG) from behavior (fine-tuning), keeping user-specific data under user control, and grounding every response in curated sources while maintaining empathetic, supportive communication.

5.2 5.2 Impact on Technology

The system advances the field of AI-assisted mental health by integrating retrieval-augmented generation with clinical safety auditing. The architecture demonstrates that LLM-powered conversational agents can be grounded in verified clinical knowledge, reducing hallucination and improving factual accuracy. The multi-tier JWT authentication, row-level security policies, and bilingual crisis detection provide a template for privacy-first healthcare applications.

The gamification engine, with its XP system, streak multipliers, badges, achievements, and challenges, provides a framework for sustaining engagement in health applications. The design targets autonomy, competence, and relatedness, which are key drivers of intrinsic motivation.

5.3 5.3 Impact on Society

The system addresses the mental health treatment gap in the MENA region, where cultural stigma, shortage of professionals, and financial barriers prevent millions from seeking help. By providing accessible, stigma-free, and scalable mental health support, the system can complement traditional face-to-face therapy and reach underserved populations.

The clinical safety auditor ensures that every response is screened for crisis indicators before delivery, with mandatory redirection to emergency resources when needed. The bilingual crisis detection (English and Arabic) addresses a gap in existing digital mental health tools that primarily serve English-speaking populations.

5.4 5.4 Ethics Considerations

The deployment of an AI-powered mental health companion raises significant ethical concerns that must be addressed transparently.

Data Privacy and Confidentiality. The system collects sensitive mental health data including chat messages, emotion labels, assessment scores, and journal entries. All data is encrypted in transit (TLS 1.3) and at rest (AES-256 via Supabase). Row-Level Security (RLS) policies enforce strict user-level data isolation: no user can access another user’s data, and even system administrators cannot view message content without explicit user consent for research purposes. Data retention follows a 90-day automatic purge policy for chat messages and a 365-day policy for assessment results, after which only anonymized aggregates are retained.

Privacy by Design for On-Device Detection. The on-device harmful content detection and visual content safety scanner operate under a strict privacy-by-design model. Raw text and screenshots are processed locally on the user’s device and immediately discarded after analysis. Only metadata alerts — such as a timestamp and threat category — are dispatched to the Flutter UI through a thread-safe event bus. No raw user text or image data leaves the device, ensuring that the expanded detection capability does not compromise user privacy.

Informed Consent. Users provide explicit informed consent during onboarding, which covers: (a) the nature of AI-generated responses, (b) the limitations of the emotion classifier, (c) data collection scope and purpose, (d) the right to export and delete all personal data, and (e) crisis detection and escalation procedures. Consent can be withdrawn at any time through the profile settings, triggering immediate data deletion.

Algorithmic Bias and Fairness. The designed BERT-based emotion classifier would be trained primarily on English-language datasets, while the current keyword-based implementation relies on English keyword lists; both introduce bias against dialectal variations and code-switching common in MENA populations. The system mitigates this by flagging low-confidence classifications (below 60%) and defaulting to a neutral therapeutic response rather than a potentially biased emotion-adaptive one. The community feature employs content moderation through user reporting, but automated bias detection in community recommendations is not yet implemented and remains a limitation.

Crisis Intervention Ethics. The bilingual crisis auditor detects life-threatening keywords in English and Arabic, but it is not infallible. When a crisis is detected (RED status), the system blocks the AI response and displays localized crisis hotline information. The system never withholds human help: users are always presented with the option to contact emergency services, regardless of the auditor’s classification. A false negative (missed crisis) is the most serious risk; the system addresses this by over-detecting potential crises and routing flagged conversations to the clinical auditor for secondary review.

Service Continuity Ethics. The service continuity protection mechanism — which resists casual uninstallation or force-stopping of the application — introduces a tension between user autonomy and safety. During acute crisis, users may impulsively remove safety tools that are essential for their protection. The mechanism is designed with transparency: users are informed during onboarding that the protection exists to preserve access to crisis resources, and a capacity-based opt-out process is available for users who can demonstrate informed decision-making. The protection does not defend against a determined technical adversary with ADB access; it is intended solely to prevent impulsive disabling during moments of vulnerability.

Therapeutic Boundaries. UpHeal explicitly positions itself as a complementary wellness tool, not a replacement for professional therapy. The onboarding flow and settings screen display a persistent disclaimer: “UpHeal is not a substitute for professional mental health care. If you are in crisis, contact emergency services immediately.” The gamification system is designed to encourage engagement, not to incentivize dependency; streak multipliers decay after 48 hours of inactivity, and no feature is locked behind XP thresholds that could pressure vulnerable users.

Transparency and Explainability. Each AI response is accompanied by an emotion badge showing the detected emotional state, giving users insight into how the system interpreted their input. The clinical auditor’s safety classification (GREEN/YELLOW/RED) is logged but not surfaced to users to avoid alarm fatigue; however, when a response is modified or blocked, the user receives a clear explanation of why.

5.5 5.5 Suggestions for Future Work

The following directions are prioritized based on feasibility, impact, and dependence on the current implementation:

1. **Full LLM integration in backend chat service.** The current prototype routes LLM calls from the Flutter frontend through the backend gateway to Groq Cloud, but the backend `ChatService._generate_response()` method returns a placeholder response. Implementing the full pipeline — prompt assembly with RAG-retrieved context, emotion label, conversation history, clinical auditing, and response storage — would enable server-side response caching, centralized cost tracking, and consistent safety auditing across all response paths.
2. **Deploy fine-tuned emotion classifier.** The prototype uses keyword-based sentiment analysis with clinical safety keywords. Fine-tuning a BERT-based classifier on a domain-specific mental health dataset — such as DAIC-WOZ or a custom Arabic-English corpus — would improve classification accuracy, particularly for the “Distressed” category (current mock $F1 = 0.71$ for the designed model). Integration with a speech emotion recognition model via Whisper [17] would enable multimodal emotion detection.
3. **Speech-to-text via Whisper.** Adding voice input reduces the barrier to engagement, particularly for users with limited literacy or those who prefer speaking over typing. Whisper’s multilingual capabilities would also support Arabic voice input.
4. **Full Arabic language support.** While the system includes Arabic crisis keyword detection and the Flutter UI supports Arabic localization via the `flutter_localizations` package, full Arabic support requires: (a) Arabic emotion classification training data, (b) Arabic clinical documents for the RAG knowledge base, (c) Arabic-tuned LLM prompting, and (d) validation with Arabic-speaking users. Given the MENA target market, this is a high-impact direction.
5. **Therapist dashboard.** A web-based clinical supervision interface would allow mental health professionals to review flagged conversations, monitor patient progress, adjust therapeutic plans, and intervene when the system detects elevated risk. This would bridge the gap between fully automated support and professional oversight.
6. **Longitudinal clinical study.** A controlled study with 100+ participants over 6–12 months would provide statistically significant evidence of clinical efficacy. Outcome measures should include PHQ-9, GAD-7, and the Working Alliance Inventory.
7. **Model quantization for edge deployment.** Quantizing the LLM and emotion classifier to 4-bit or 8-bit precision would enable on-device inference, reducing latency and eliminating cloud dependency for privacy-sensitive deployments. QLoRA [13] provides an efficient fine-tuning and quantization framework for this purpose. This is particularly relevant for the planned self-hosted Ollama migration [26].
8. **Multi-agent conversation architecture.** Future iterations could employ a multi-agent system where separate specialized agents handle crisis detection, therapeutic intervention, gamification coaching, and assessment delivery, coordinated by a routing agent. This modular design would improve safety isolation and enable independent scaling of each agent.
9. **Cloudflare edge protection for production.** Deploying Cloudflare WAF, DDoS mitigation, bot management, and Turnstile attestation at the DNS edge would harden the platform against automated attacks and credential stuffing before traffic reaches the origin servers.
10. **Certificate pinning for mobile clients.** Pinning the backend TLS certificate inside the Flutter

application would eliminate certain man-in-the-middle attack vectors on untrusted networks, complementing TLS 1.3 and perfect forward secrecy.

11. **mTLS for inter-service communication.** If the architecture evolves from the current in-process gateway to distributed microservices, mutual TLS should be enforced for all internal service-to-service calls to maintain Zero Trust verification between components.
12. **Biometric authentication.** Integrating fingerprint or face recognition as a second authentication factor would strengthen account security while preserving usability, particularly for users who may forget passwords during periods of cognitive distress.

5.6 5.6 Limitations

The current implementation has several limitations that should be acknowledged:

1. **Prototype-grade emotion classifier.** The emotion classifier has not been fine-tuned on a mental health-specific dataset and may underperform on clinical-grade emotion detection tasks.
2. **Small evaluation cohort.** The user testing cohort (n=50 beta testers, 200 annotated responses) is insufficient for statistically robust conclusions. Larger-scale studies are needed to validate the findings.
3. **Limited Arabic support.** While the system includes bilingual crisis detection, the core NLP pipeline — emotion classification (currently keyword-based sentiment), RAG retrieval, and LLM prompting — operates primarily in English, limiting applicability for Arabic-speaking users.
4. **Cloud dependency.** The current architecture depends on Groq Cloud for LLM inference and Supabase for data persistence. A fully self-hosted deployment would require significant infrastructure investment.
5. **No clinical validation.** The system has not undergone clinical trials and should not be considered a replacement for professional mental health care. It is designed as a complementary tool.
6. **Community moderation.** The community feature relies on user reports for content moderation; automated content detection is not yet implemented.
7. **On-device detection accuracy.** The on-device harmful content detector relies on keyword matching and a pre-trained TensorFlow Lite toxicity model. Fine-tuning on mental health-specific content datasets would improve accuracy and reduce false positives for clinical language.
8. **Visual scanner overhead.** The visual content safety scanner processes screenshots every four seconds, which introduces CPU and battery overhead. Adaptive interval strategies that scale with user risk state should be explored.
9. **Anti-tampering limitations.** The service continuity mechanism prevents casual uninstallation and force-stopping, but it does not protect against a determined technical user with ADB access or custom recovery tools.

5.7 5.7 Conclusion

UpHeal demonstrates that a fine-tuned, RAG-grounded LLM combined with structured memory, emotion-aware context processing, and real-time voice interaction can deliver more personalized, empathetic, and safer mental-health support than standalone chatbots. The system provides a foundation for scalable, clinically grounded, and culturally adaptable digital mental-health assistance in the MENA region and beyond.

Appendix A

Appendix

A.1 A.1 Source Code Repositories

The complete source code is available at the following GitHub repositories under the MIT License:

Repository	Description	URL
UpHeal-Doc	Thesis documentation and Quarto project	https://github.com/Upheal1/UpHeal-Doc
UpHeal-Backend	Python FastAPI backend and AI pipeline	https://github.com/Upheal1/UpHeal-Backend
UpHeal-Frontend	Flutter mobile application	https://github.com/Upheal1/UpHeal-Frontend

A.2 A.2 Publications

No publications have resulted from this project at the time of submission.

- [1] World Health Organization, “Mental health: Strengthening our response.” <https://www.who.int/news-room/fact-sheets/detail/mental-health-strengthening-our-response>, 2021.
- [2] P. Lewis *et al.*, “Retrieval-augmented generation for knowledge-intensive NLP tasks,” *Advances in Neural Information Processing Systems*, vol. 33, pp. 9459–9474, 2020.
- [3] J. Hamari, J. Koivisto, and H. Sarsa, “Gamification: A systematic literature review of gamification research in the context of health and well-being,” 2016.
- [4] Y. Chen, B. Vargas, P. Dutta, and P. Rui, “Mobile health app usage and retention: A systematic review,” *Journal of Medical Internet Research*, vol. 22, no. 7, p. e18135, 2020.
- [5] A. Vaswani *et al.*, “Attention is all you need,” *Advances in Neural Information Processing Systems*, vol. 30, 2017.
- [6] J. Devlin, M.-W. Chang, K. Lee, and K. Toutanova, “BERT: Pre-training of deep bidirectional transformers for language understanding,” *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics*, pp. 4171–4186, 2019.
- [7] L. Ouyang *et al.*, “Training language models to follow instructions with human feedback,” *Advances in Neural Information Processing Systems*, vol. 35, pp. 27730–27744, 2022.
- [8] R. Fulmer, A. Joerin, B. Gentile, L. Loughren, and E. Ralston, “Woebot: A therapeutic chatbot for depression and anxiety,” in *Proceedings of the ACM human-computer interaction*, ACM, 2021, pp. 1–19.
- [9] B. Inkster, S. Sarda, and V. Subramanian, “An empathy-driven, conversational artificial intelligence agent (wysa) for digital mental well-being,” *JMIR mHealth and uHealth*, vol. 6, no. 11, p. e12106, 2018.

- [10] A. Kumar and A. Jaiswal, “Mental health conversational agents: A systematic review of architectures, datasets, and evaluation metrics,” *Computers in Human Behavior*, vol. 148, p. 107910, 2023.
- [11] H. Rashkin, E. M. Smith, M. Li, and Y.-L. Boureau, “Towards empathetic open-domain conversation models: A new benchmark and dataset,” in *Proceedings of the 27th international conference on computational linguistics*, 2018, pp. 537–548.
- [12] E. J. Hu *et al.*, “LoRA: Low-rank adaptation of large language models,” *arXiv preprint arXiv:2106.09685*, 2021.
- [13] T. Dettmers, A. Pagnoni, A. Holtzman, and L. Zettlemoyer, “QLoRA: Efficient fine-tuning of quantized large language models,” *arXiv preprint arXiv:2305.14314*, 2023.
- [14] C. Packer, V. Fang, S. G. Patil, K. Lin, J. Wooders, and J. E. Gonzalez, “MemGPT: Towards LLMs as operating systems,” *arXiv preprint arXiv:2310.08560*, 2023.
- [15] D. Demszky, D. Movshovitz-Attias, J. Ko, A. Cowen, G. Nemade, and S. Ravi, “GoEmotions: A dataset of fine-grained emotions,” in *Proceedings of the 58th annual meeting of the association for computational linguistics*, 2020, pp. 4040–4054.
- [16] A. Baevski, Y. Zhou, A. Mohamed, and M. Auli, “wav2vec 2.0: A framework for self-supervised learning of speech representations,” *Advances in Neural Information Processing Systems*, vol. 33, pp. 12449–12460, 2020.
- [17] A. Radford, J. W. Kim, T. Xu, G. Brockman, C. McLeavey, and I. Sutskever, “Robust speech recognition via large-scale weak supervision,” *arXiv preprint arXiv:2212.04356*, 2022.
- [18] E. Casanova *et al.*, “XTTS: A massively multilingual zero-shot text-to-speech model,” *arXiv preprint arXiv:2406.04946*, 2024.
- [19] Y. Chen *et al.*, “F5-TTS: A fairytaler that fakes fluent and faithful speech with flow matching,” *arXiv preprint arXiv:2410.06885*, 2024.
- [20] Z. Guo, A. Lai, J. H. Thygesen, J. Farrington, T. Keen, and K. Li, “Large language model for mental health: A systematic review,” *arXiv preprint arXiv:2401.00735*, 2024.
- [21] Y. Hua *et al.*, “Large language models in mental health care: A scoping review,” *arXiv preprint arXiv:2401.00286*, 2024.
- [22] World Health Organization, “Ethics and governance of artificial intelligence for health: WHO guidance,” World Health Organization; <https://www.who.int/publications/i/item/9789240029200>, 2021.
- [23] Gemma Team, “Gemma 3 technical report,” Google DeepMind; <https://storage.googleapis.com/deepmind-media/gemma/Gemma3Report.pdf>, 2025.
- [24] D. Seo, G. Park, and K. Nam, “Chatterbox-flash: Prior-calibrated block diffusion for streaming zero-shot TTS,” *arXiv preprint arXiv:2502.05538*, 2026.
- [25] Qdrant, “Qdrant: Vector search engine.” <https://qdrant.tech/>, 2024.
- [26] Gemma Team, “Gemma: Open models based on Google research and technology,” *arXiv preprint arXiv:2403.08295*, 2024.